

Lesson 7

PRACTICAL DEEP LEARNING FOR CODERS 2022

WHAT'S INSIDE A NEURAL NET?

Entity Embeddings of Categorical Variables

Cheng Guo^{*} and Felix Berkhahn[†]

Neokami Inc.

(Dated: April 25, 2016)

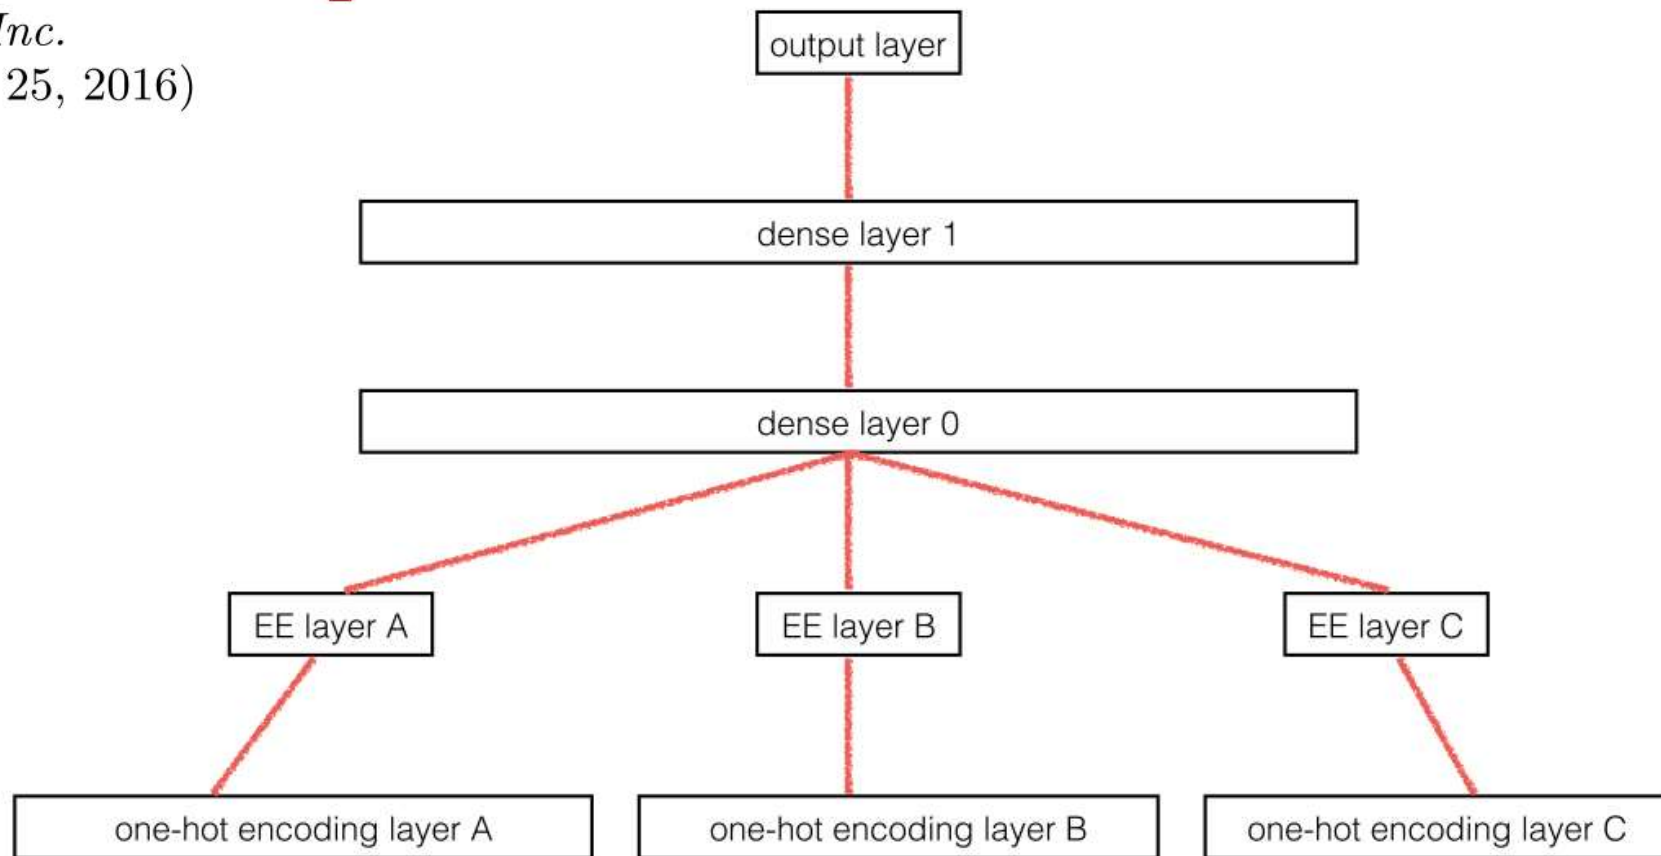
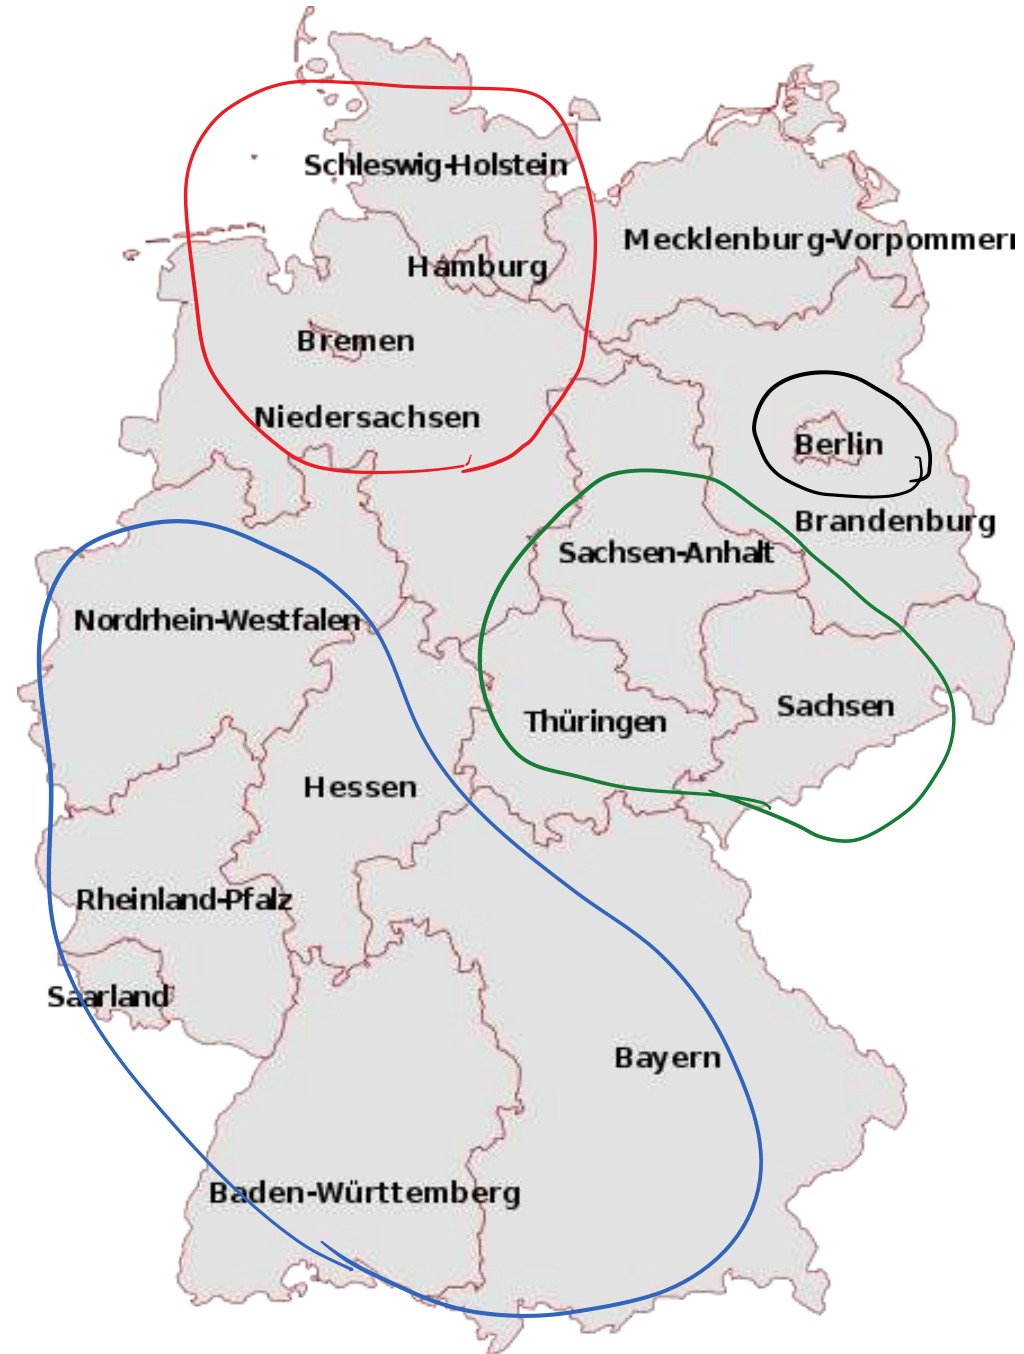
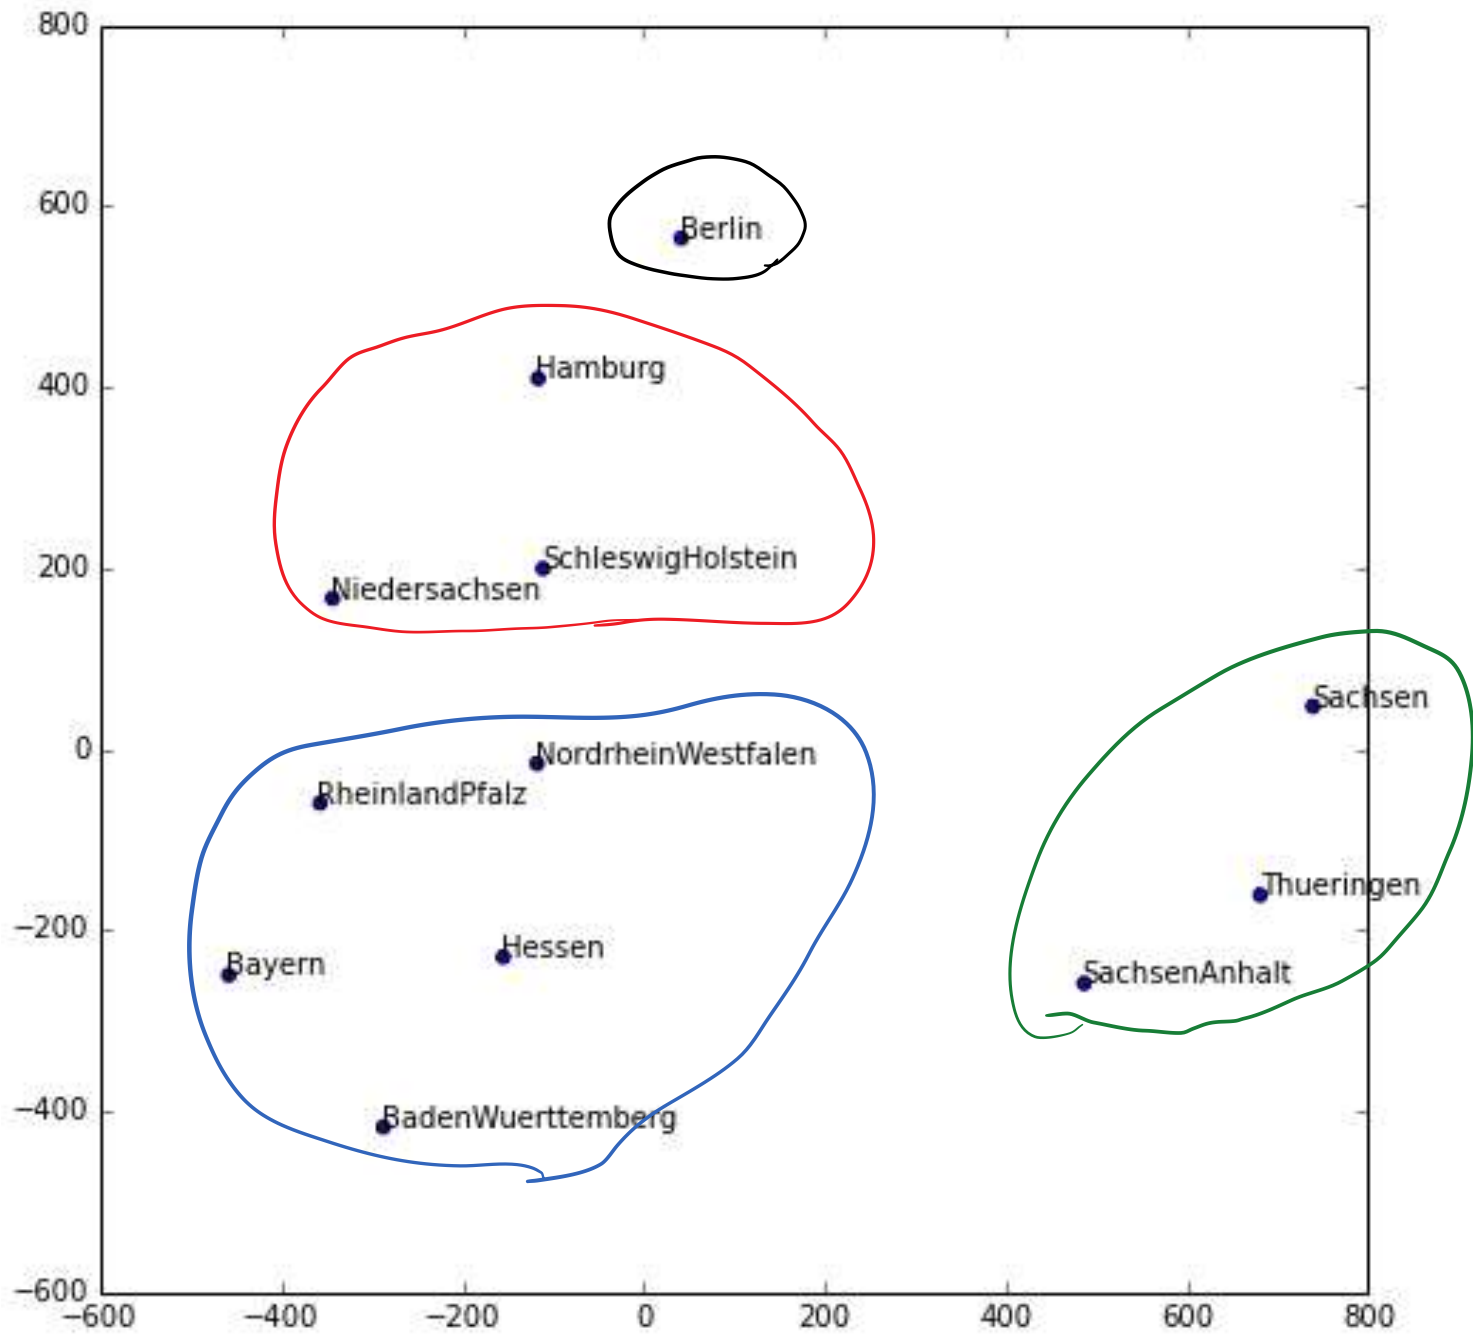
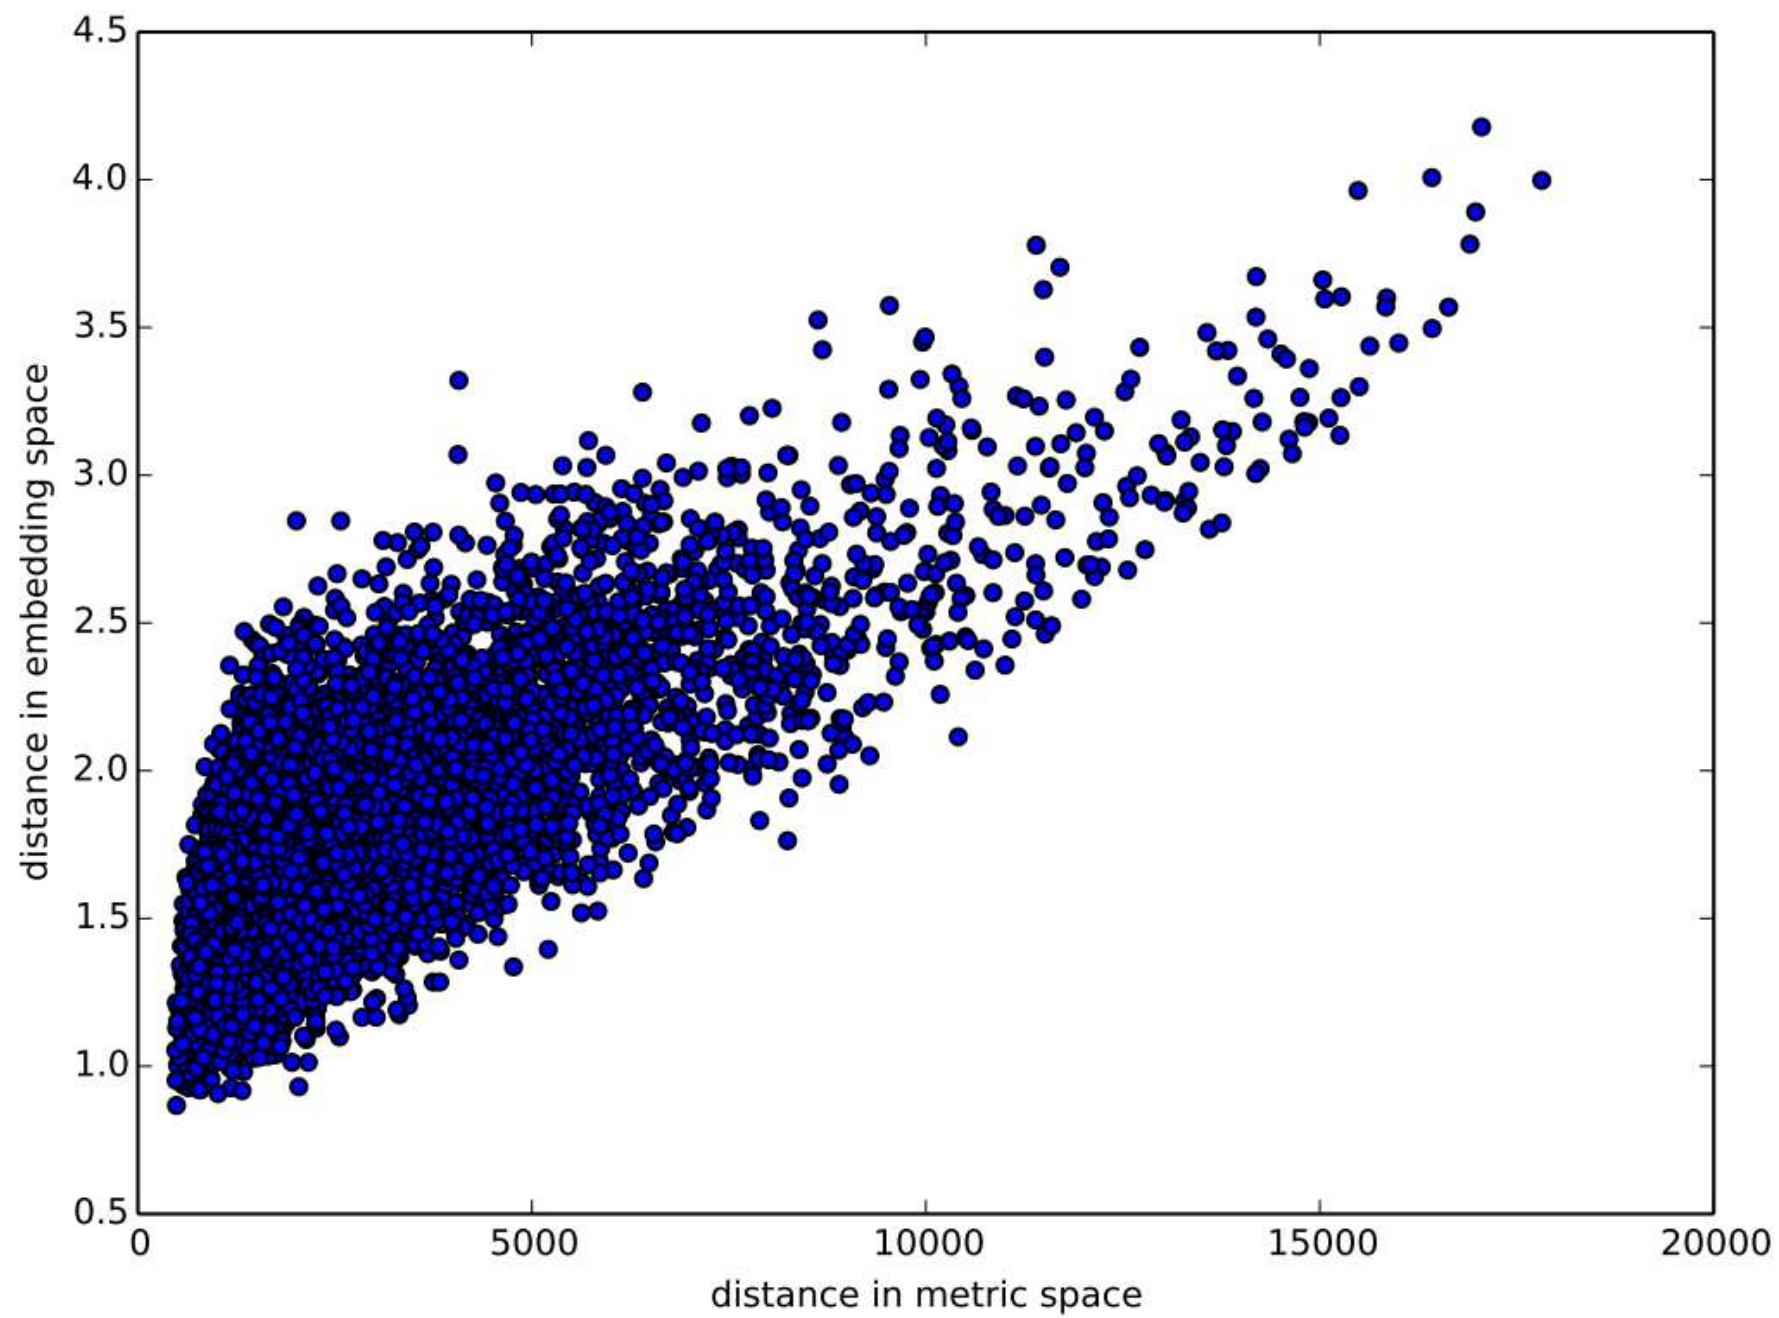
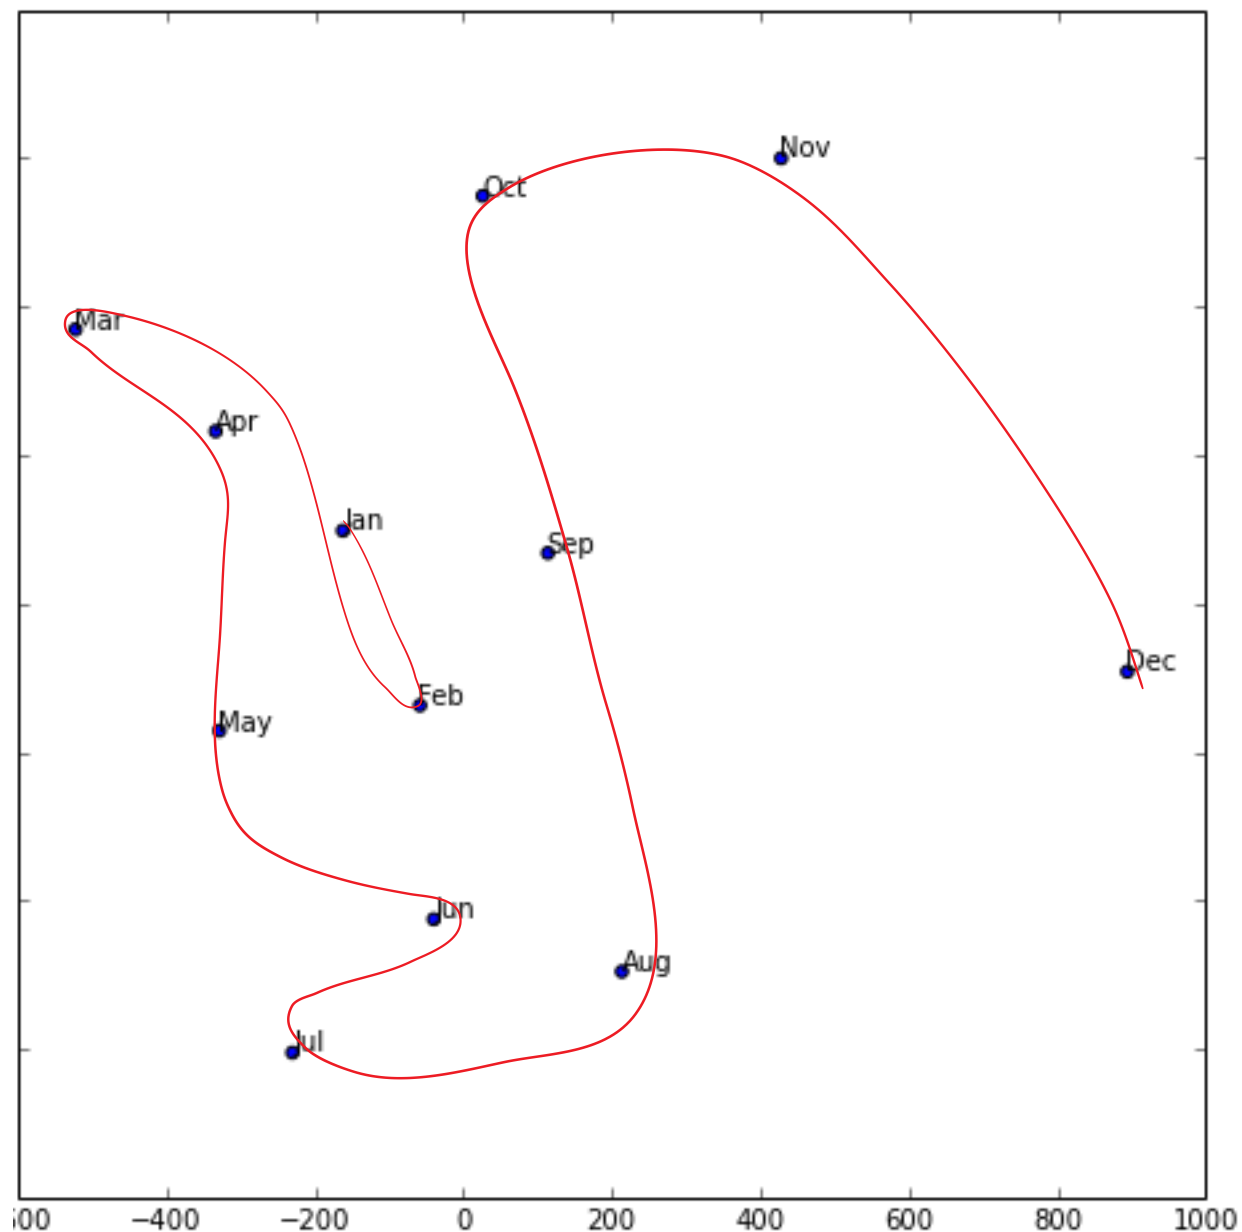
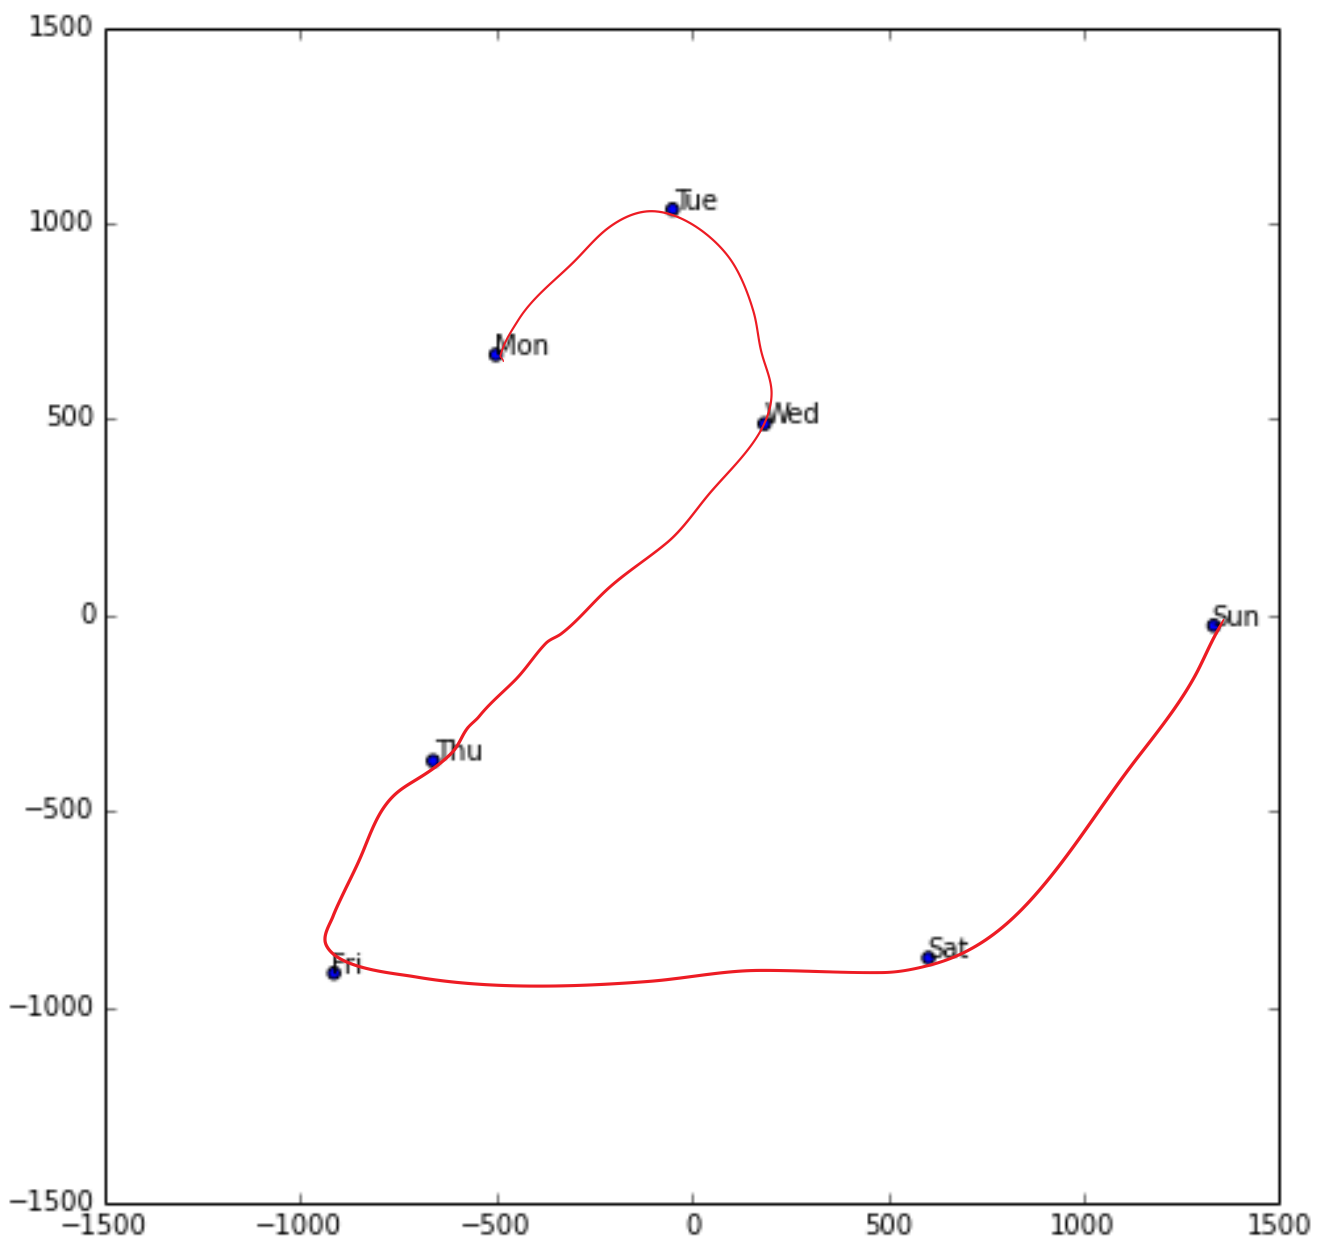


FIG. 1. Illustration that entity embedding layers are equivalent to extra layers on top of each one-hot encoded input.

method	MAPE	MAPE (with EE)
KNN	0.290	0.116
random forest	0.158	0.108
gradient boosted trees	0.152	0.115
neural network	0.101	0.093









Matthew Kleinsmith [Follow](#)

Studying deep learning

Feb 10, 2017 · 4 min read

CNNs from different viewpoints

Prerequisite: Basic neural networks

$$\alpha A + \beta B + \gamma D + \delta E + b = P$$

$$\alpha B + \beta C + \gamma E + \delta F + b = Q$$

$$\alpha D + \beta E + \gamma G + \delta H + b = R$$

$$\alpha E + \beta F + \gamma H + \delta J + b = S$$

α	β
γ	δ

applied to

A	B	C
D	E	F
G	H	J

yields

P	

α	β
γ	δ

A	B	C
D	E	F
G	H	J

	Q

α	β
γ	δ

A	B	C
D	E	F
G	H	J

R	

α	β
γ	δ

A	B	C
D	E	F
G	H	J

	S

α	β	0	γ	δ	0	0	0	0
0	α	β	0	γ	δ	0	0	0
0	0	0	α	β	0	γ	δ	0
0	0	0	0	α	β	0	γ	δ

*

A
B
C
D
E
F
G
H
J

+

b
b
b
b

=

$\alpha A + \beta B + 0C + \gamma D + \delta E + 0F + 0G + 0H + 0J + b$
$0A + \alpha B + \beta C + 0D + \gamma E + \delta F + 0G + 0H + 0J + b$
$0A + 0B + 0C + \alpha D + \beta E + 0F + \gamma G + \delta H + 0J + b$
$0A + 0B + 0C + 0D + \alpha E + \beta F + 0G + \gamma H + \delta J + b$

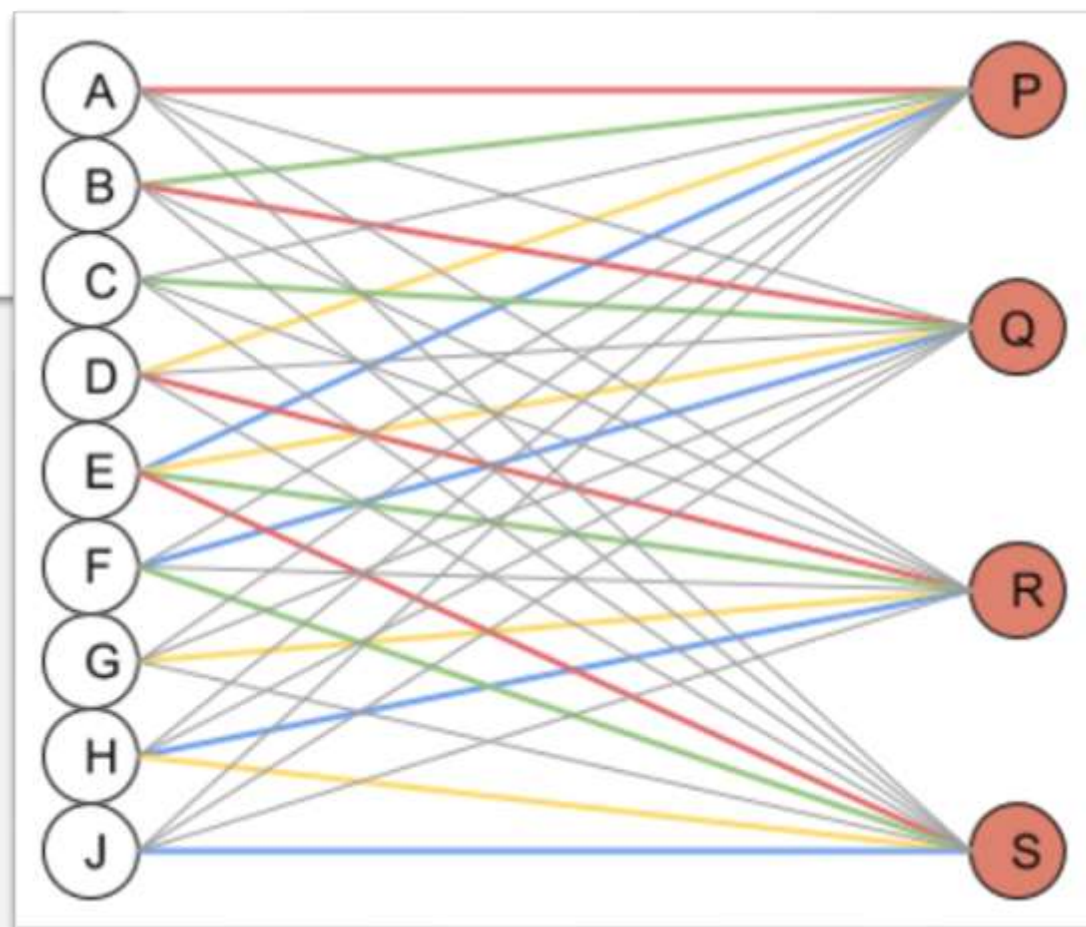
=

$\alpha A + \beta B + \gamma D + \delta E + b$
$\alpha B + \beta C + \gamma E + \delta F + b$
$\alpha D + \beta E + \gamma G + \delta H + b$
$\alpha E + \beta F + \gamma H + \delta J + b$

=

P
Q
R
S

A B C D E F G H J



α	β	γ
δ	ϵ	ζ
η	θ	κ

applied to

0	0	0	0	0
0	A	B	C	0
0	D	E	F	0
0	G	H	J	0
0	0	0	0	0

yields

A'		


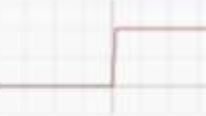







α	β	γ
δ	ϵ	ζ
η	θ	κ

applied to

0	0	0	0	0
0	A	B	C	0
0	D	E	F	0
0	G	H	J	0
0	0	0	0	0

yields

	B'	

Name	Plot	Equation	Derivative
Identity		$f(x) = x$	$f'(x) = 1$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$
Logistic (a.k.a Soft step)		$f(x) = \frac{1}{1 + e^{-x}}$	$f'(x) = f(x)(1 - f(x))$
TanH		$f(x) = \tanh(x) = \frac{2}{1 + e^{-2x}} - 1$	$f'(x) = 1 - f(x)^2$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$
Rectified Linear Unit (ReLU)		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$
Parameteric Rectified Linear Unit (PReLU) [2]		$f(x) = \begin{cases} \alpha x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	
Exponential Linear Unit (ELU) [3]		$f(x) = \begin{cases} \alpha(e^x - 1) & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	
SoftPlus		$f(x) = \log_e(1 + e^x)$	



SAGAR SHARMA

Follow

I am interested in Programming (Python, C++), Arduino, Machine learning :) I'm the editor of Arduino Community on Medium. I also like to write stuff.

Sep 6, 2017 · 5 min read

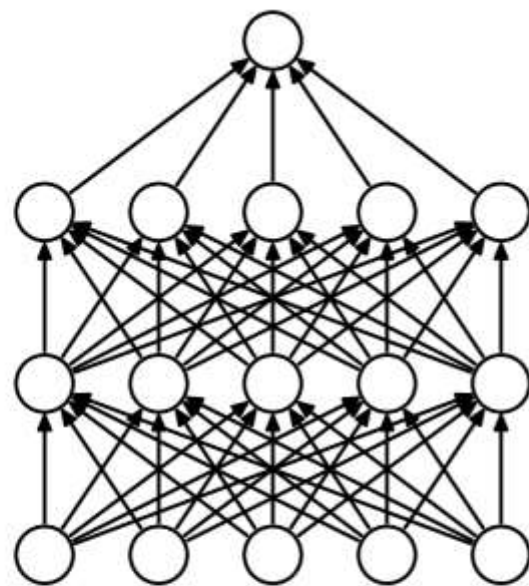
Activation Functions: Neural Networks

Sigmoid, tanh, Softmax, ReLU, Leaky ReLU EXPLAINED !!!

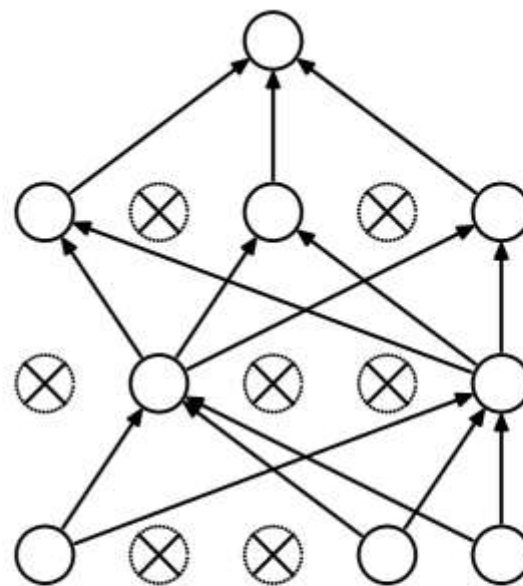
Dropout: A Simple Way to Prevent Neural Networks from Overfitting

Nitish Srivastava
Geoffrey Hinton
Alex Krizhevsky
Ilya Sutskever
Ruslan Salakhutdinov

NITISH@CS.TORONTO.EDU
HINTON@CS.TORONTO.EDU
KRIZ@CS.TORONTO.EDU
ILYA@CS.TORONTO.EDU
RSALAKHU@CS.TORONTO.EDU



(a) Standard Neural Net



(b) After applying dropout.

Figure 1: Dropout Neural Net Model. **Left:** A standard neural net with 2 hidden layers. **Right:** An example of a thinned net produced by applying dropout to the network on the left. Crossed units have been dropped.

“I went to my bank. The tellers kept changing and I asked one of them why. He said he didn’t know but they got moved around a lot. I figured it must be because it would require cooperation between employees to successfully defraud the bank. This made me realize that randomly removing a different subset of neurons on each example would prevent conspiracies and thus reduce overfitting”

Hinton: Reddit AMA



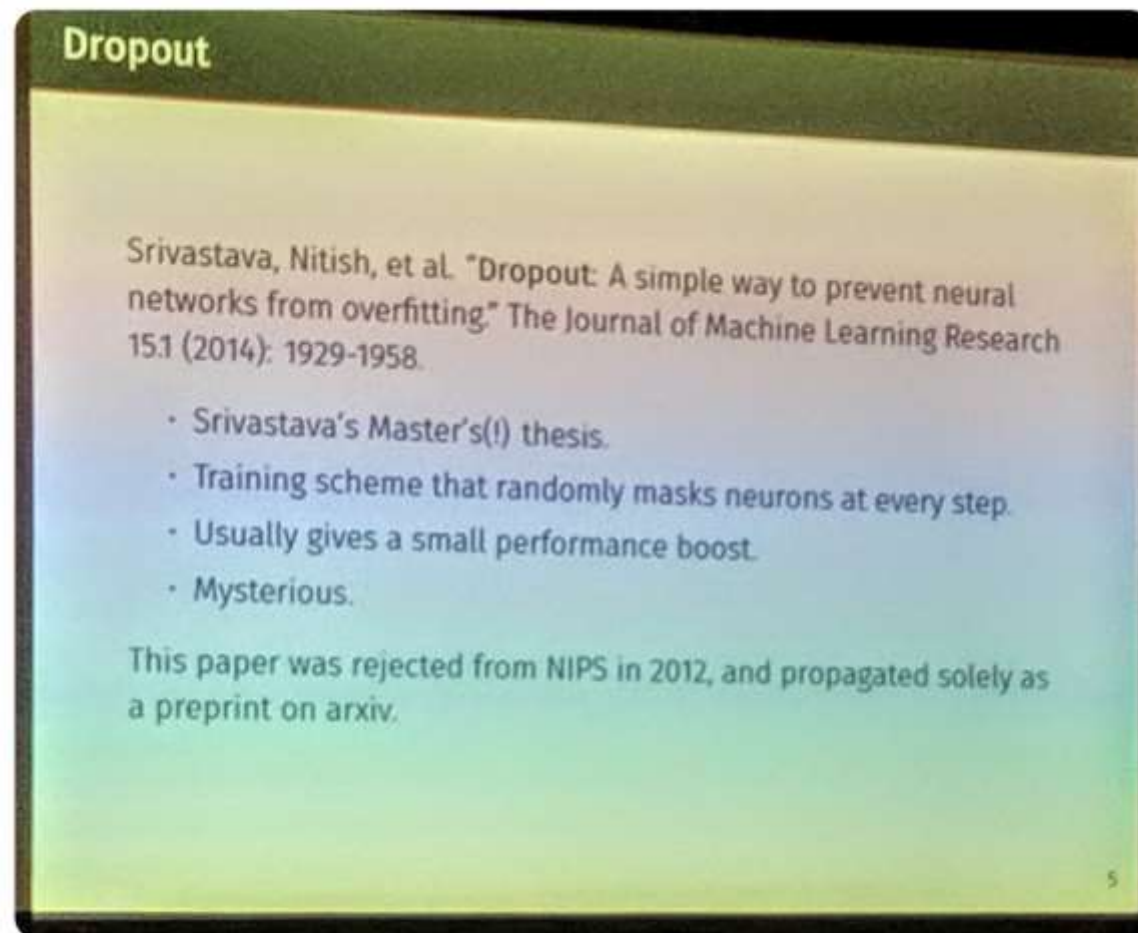
Chris Gorgolewski

@ChrisFiloG

Follow



Did you know that Dropout was originally introduced in a Master's thesis and was rejected from NIPS? Was disseminated via #arxiv! #OHBM2018



Dropout

Srivastava, Nitish, et al. "Dropout: A simple way to prevent neural networks from overfitting." *The Journal of Machine Learning Research* 15.1 (2014): 1929-1958.

- Srivastava's Master's(!) thesis.
- Training scheme that randomly masks neurons at every step.
- Usually gives a small performance boost.
- Mysterious.

This paper was rejected from NIPS in 2012, and propagated solely as a preprint on arxiv.

5

So what now? Watch the videos again, and...

